

Advanced Image Preprocessing and Augmentation for Deep Learning (DL)

Assignment Objective

The goal of this assignment is to implement a robust image preprocessing and augmentation pipeline that mimics the workflow necessary to prepare raw images for training a Convolutional Neural Network (CNN). You will explore fundamental techniques across filtering, color space transformation, noise management, and data augmentation.

Required Setup and Tools

- **Programming Language:** Python 3.x
- **Libraries:**
 - `numpy` (for array manipulation)
 - `opencv-python (cv2)` or `Pillow (PIL)` (for core image operations)
 - `matplotlib` (for visualization)
- **Input Data:** Use a sample set of 5–10 diverse images (e.g., common images like flowers, animals, or general scenes).

Task 1: Essential Preprocessing and Normalization

These steps ensure all input images conform to the requirements of a neural network model architecture (e.g., VGG, ResNet).

1. **Uniform Resizing:**
 - Implement a function that takes an input image and resizes it to a fixed dimension, $W \times H$. Use 224×224 pixels as the target size.
 - Ensure the function can handle images with different initial aspect ratios (you may choose to stretch the image or pad it, but stretching is often the simplest initial approach).
2. **Color Channel Format Conversion:**
 - Verify and, if necessary, convert the image color ordering to the standard format required by most DL frameworks (RGB: Red, Green, Blue). Note: OpenCV often uses BGR by default, so conversion might be necessary.
3. **Data Type and Normalization:**

- Convert the image array data type to floating-point (e.g., `float32`).
- Implement **Min-Max Normalization**: Scale the pixel values from the standard range to the floating-point range of [0.0,1.0].
- *Challenge (Optional)*: Implement **Standardization** (Z-score normalization) using pre-calculated global mean (μ) and standard deviation (σ) across all training data (for this assignment, you can calculate the μ and σ of your 5–10 sample images).

Task 2: Filtering and Noise Management

Noise is common in real-world data. Filtering is critical for both removing noise and highlighting certain image features.

1. Noise Addition:

- Implement a function to deliberately corrupt an image by adding **Salt-and-Pepper Noise**.
- Implement a function to corrupt an image by adding **Gaussian Noise** (additive white noise).
- Apply both noise types to the same original image, creating two distinct corrupted versions.

2. Noise Removal (Denoising):

- Apply a **Median Filter** (e.g., 3×3 or 5×5 kernel) to the **Salt-and-Pepper Noise** corrupted image.
- Apply a **Gaussian Blur/Filter** (e.g., $\sigma=1.0$ or 2.0) to the **Gaussian Noise** corrupted image.

3. Quantitative Evaluation:

- Calculate the **Peak Signal-to-Noise Ratio (PSNR)** and **Mean Squared Error (MSE)** between: a) The original image and the noisy image. b) The original image and the denoised image (after filtering).
- Compare the metrics to demonstrate the effectiveness of the chosen filters.

Task 3: Color Transformation and Data Augmentation

Data augmentation is a crucial preprocessing technique to prevent overfitting in DL models.

1. Color Space Transformation:

- Implement a function to convert the image from **RGB to Grayscale**. (Grayscale images are often used in certain CNN tasks, or as a channel for color augmentation).

2. Hue and Saturation Adjustment:

- Convert an image from **RGB to the HSV (Hue, Saturation, Value) color space.**
- Implement an operation to **randomly adjust the Saturation** of the image by a factor between 0.5 and 1.5.
- Convert the image back from HSV to RGB for display.

3. Geometric Augmentation Pipeline:

- Create a single augmentation function that randomly applies the following techniques with a 50% probability for each: a) **Horizontal Flip:** Mirror the image along the vertical axis. b) **Random Rotation:** Rotate the image by a small random angle between -10° and $+10^\circ$. Fill any resulting empty pixels with a neutral color (e.g., the image mean).

Deliverables

1. **Python Script:** A well-commented Python file (`preprocessing_pipeline.py`) containing all implemented functions and a demonstration of each task.
2. **Visual Report:** A set of visualizations (using `matplotlib`) showing the results for at least one sample image, including:
 - Original Image.
 - Resized and Normalized Image.
 - Image with Salt-and-Pepper Noise and the Median Filtered Result (with PSNR/MSE comparison).
 - Image with Gaussian Noise and the Gaussian Filtered Result (with PSNR/MSE comparison).
 - Image after Saturation Adjustment.
 - Image after Random Geometric Augmentation.